

# CSCI320 – Operating Systems

Ahmad Fadlallah

# Introduction

# References

- These slides are based on the official slides of the following textbooks
  - **Operating Systems: Internals and Design Principles (9<sup>th</sup> Edition)**, *William Stallings*, Pearson
  - **Operating System Concepts with Java (10<sup>th</sup> Edition)**, *Abraham Silberschatz, Peter B. Galvin and Greg Gagne*, Wiley

# Objectives

- To describe the basic organization of computer systems
- To provide a grand tour of the major components of operating systems
- To give an overview of the many types of computing environments

# What Does the Term Operating System Mean?

- An operating system is “\_\_\_\_\_”
- What about:
  - Car
  - Airplane
  - Printer
  - Washing Machine
  - Toaster
  - Compiler
  - Etc.

# What is an Operating System?

- A **program** that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals
  - **Execute user programs** and make solving user problems easier
  - Make the computer system **convenient** to use
  - Use the computer hardware in an **efficient manner**



# Computer System Structure

## ■ Components of a Computer System

### ○ Hardware

- Provides basic computing resources
  - CPU, memory, I/O devices

### ○ Operating system


- Controls and coordinates use of hardware *among various applications and users*

### ○ Application programs

- Define the ways in which the system resources are used to solve the computing problems of the users
  - Word processors, compilers, web browsers, database systems, video games

### ○ Users

- People, machines, other computers



# What Operating Systems Do?

- Depends on the point of view
- **Users** want *convenience, ease of use* and *good performance*
  - **Don't care** about *resource utilization*
- But **shared computer** (mainframe, minicomputer) **must keep all users happy**
- **Users of dedicated systems** (e.g., workstations) have dedicated resources but frequently use shared resources from servers
- **Mobile devices** (smartphones, tablets) are **resource poor, optimized for usability and battery life**
  - **Mobile user interfaces** such as touch screens, voice recognition
- **Some computers have little or no user interface**, such as embedded computers in devices and automobiles
  - **Run primarily without user intervention**



# What Operating Systems Do?

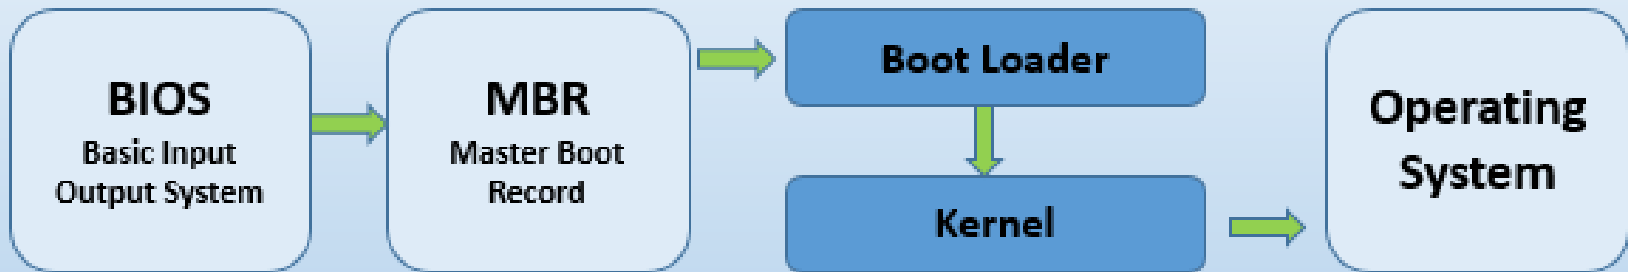
- OS is a **Resource Allocator**
  - Manages all *resources*
  - Decides between *conflicting requests* for **efficient and fair resource use**
- OS is a **Control Program**
  - Controls execution of programs to **prevent errors** and **improper use** of the computer

# Computer Startup


## BIOS vs UEFI

A Very Brief Explanation

### BIOS BOOT



### UEFI BOOT



**UEFI:** Unified Extensible Firmware Interface

**BIOS:** Basic Input/Output System

# Computer Startup

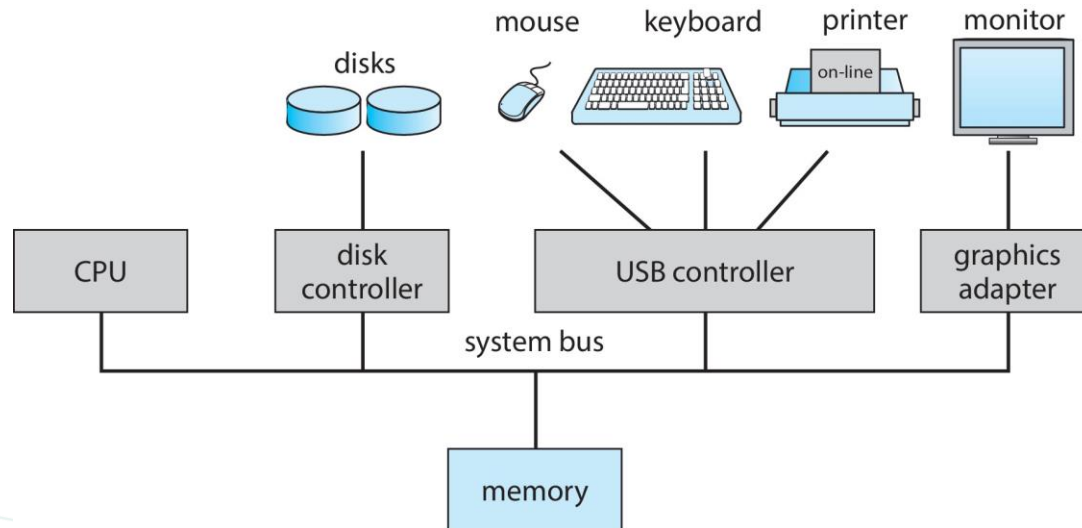
- **Bootstrap program** is loaded at power-up or reboot
  - Typically stored in *ROM/EPROM/EEPROM*, generally known as **firmware**
  - Initializes all aspects of system from *CPU registers* to *device controllers* to *memory contents*
  - Loads **operating system kernel** and starts execution
    - the bootstrap program must locate the operating-system kernel and load it into memory.

# Computer Startup (cont'd)

- Once the kernel is loaded and executing, it can start providing services to the system and its users.
- Some services are provided outside of the kernel, by **system programs** that are loaded into memory at boot time to become **system processes**, or **system daemons** that run the entire time the kernel is running.
- Once this phase is complete, the system is fully booted, and the system **waits for some event** to occur.

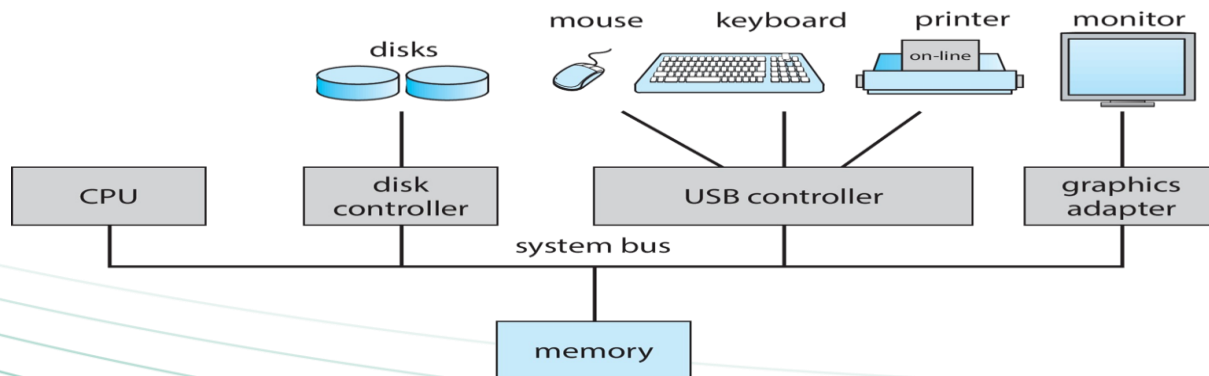
# Computer System Organization

- *One or more CPUs, device controllers connect through common bus providing access to shared memory*
- ***Concurrent*** execution of CPU(s) and devices *competing* for ***memory cycles***




# Computer-System Operation

- I/O devices and the CPU can execute **concurrently**
- Each **device controller** is in charge of a particular device type
- Each device controller has a **local buffer**
- Each device controller type has an operating system **device driver** to manage it
- **CPU moves data from/to main memory to/from local buffers**
- I/O is from the device to local buffer of controller
- **Device controller informs CPU** that it has finished its operation by causing an **interrupt**



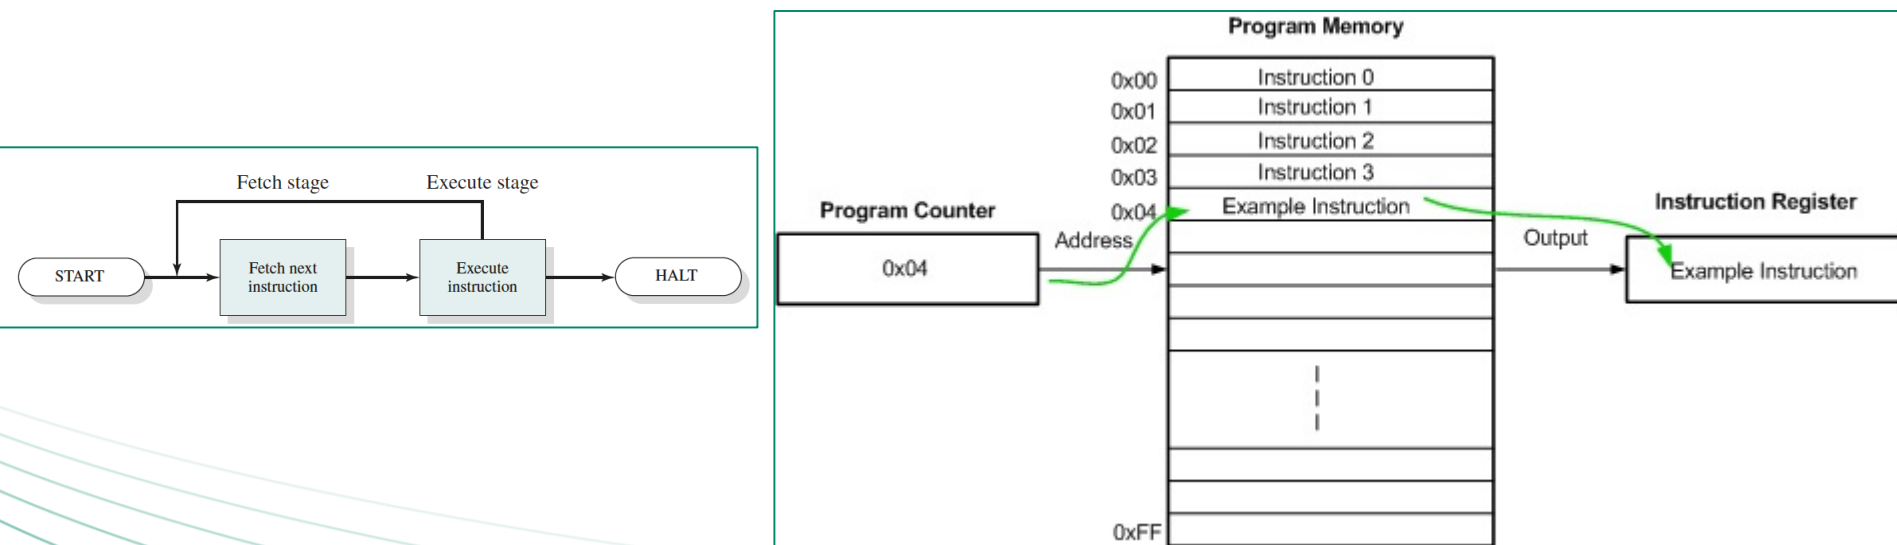
# Instruction Execution

- A **program** to be executed by a processor consists of a **set of instructions** stored in memory.
- In its simplest form, instruction processing consists of two steps:
  - The processor reads (**fetches**) instructions from memory one at a time and **executes** each instruction.
- Program execution consists of repeating the process of **instruction fetch** and **instruction execution**.
  - Instruction execution may involve several operations and depends on the nature of the instruction.
- **Program execution halts** only if the processor is turned off, some sort of unrecoverable error occurs, or a program instruction that halts the processor is encountered.



# Instruction Execution (cont'd)

- At the beginning of each instruction cycle, the processor fetches an instruction from memory.
- **The Program Counter (PC)** holds the address of the next instruction to be fetched.
  - Unless instructed otherwise, the processor always increments the PC after each instruction fetch
- The **fetch instruction** is loaded into the **Instruction register (IR)**.






# Instruction Execution (cont'd)

## ■ Instruction Action Categories

- **Processor-memory:** Data may be transferred from processor to memory, or from memory to processor.
- **Processor-I/O:** Data may be transferred to or from a peripheral device by transferring between the processor and an I/O module.
- **Data processing:** The processor may perform some arithmetic or logic operation on data.
- **Control:** An instruction may specify that the sequence of execution be altered.
  - Example: Assigning a new address to fetch the instruction from



# Interrupts

- The occurrence of an event is usually signaled by an **interrupt** from either the *hardware* or the *software*.
  - **Hardware** may trigger an interrupt at any time by sending a signal to the CPU, usually by way of the *system bus*.
  - **Software** may trigger an interrupt by executing a special operation called a **system call** (also called a **monitor call**).

# Interrupts – Main Classes

## Program

- **Generated by some condition that occurs as a result of an instruction execution**
- Arithmetic overflow, division by zero, reference outside user's allowed memory space.

## Timer

- **Generated by a timer within the processor.**
- This allows the operating system to perform certain functions on a regular basis.

## I/O

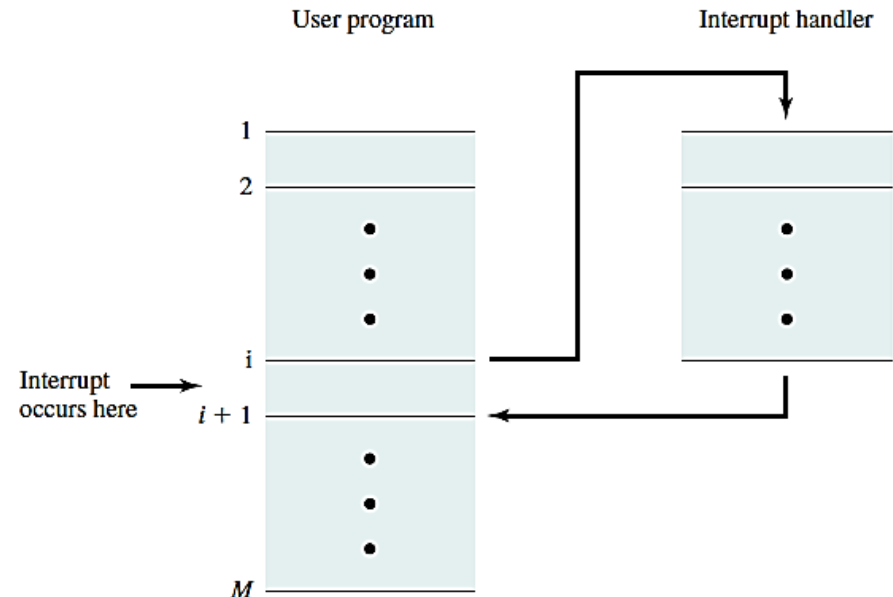
- **Generated by an I/O controller**, to signal normal completion of an operation or to signal a variety of error conditions.

## Hardware failure

- **Generated by a failure**, such as power failure or memory parity error.

# Interrupts

- *When the CPU is interrupted, it stops what it is doing and immediately transfers execution to a fixed location.*
  - The fixed location usually contains the **starting address** where the **service routine** for the interrupt (**Interrupt handler**) is located.
  - The interrupt service routine executes
  - On completion, the CPU **resumes the interrupted computation.**



# Interrupts (cont'd)

- **Why interrupts are needed?**

- Interrupts are provided primarily as a way to improve processor utilization.
- *Most I/O devices are much slower than the processor.*

- **Example1: Writing Data on a HDD**

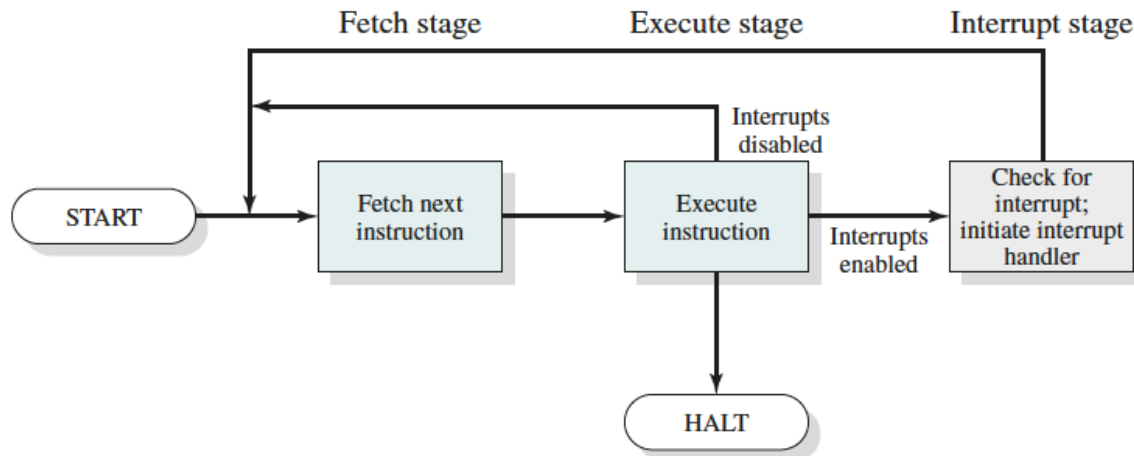

- A 1 GHz Processor is 4 millions times faster than a typical HDD (7200 RPM)

- **Example2: Printing Data**

- A processor is transferring data to a printer using the instruction cycle scheme (Fetch-Execute).
  - ✓ After each write operation, the processor must pause and remain idle until the printer catches up.
  - ✓ The length of this pause may be on the order of many thousands or even millions of instruction cycles => very wasteful use of the processor.

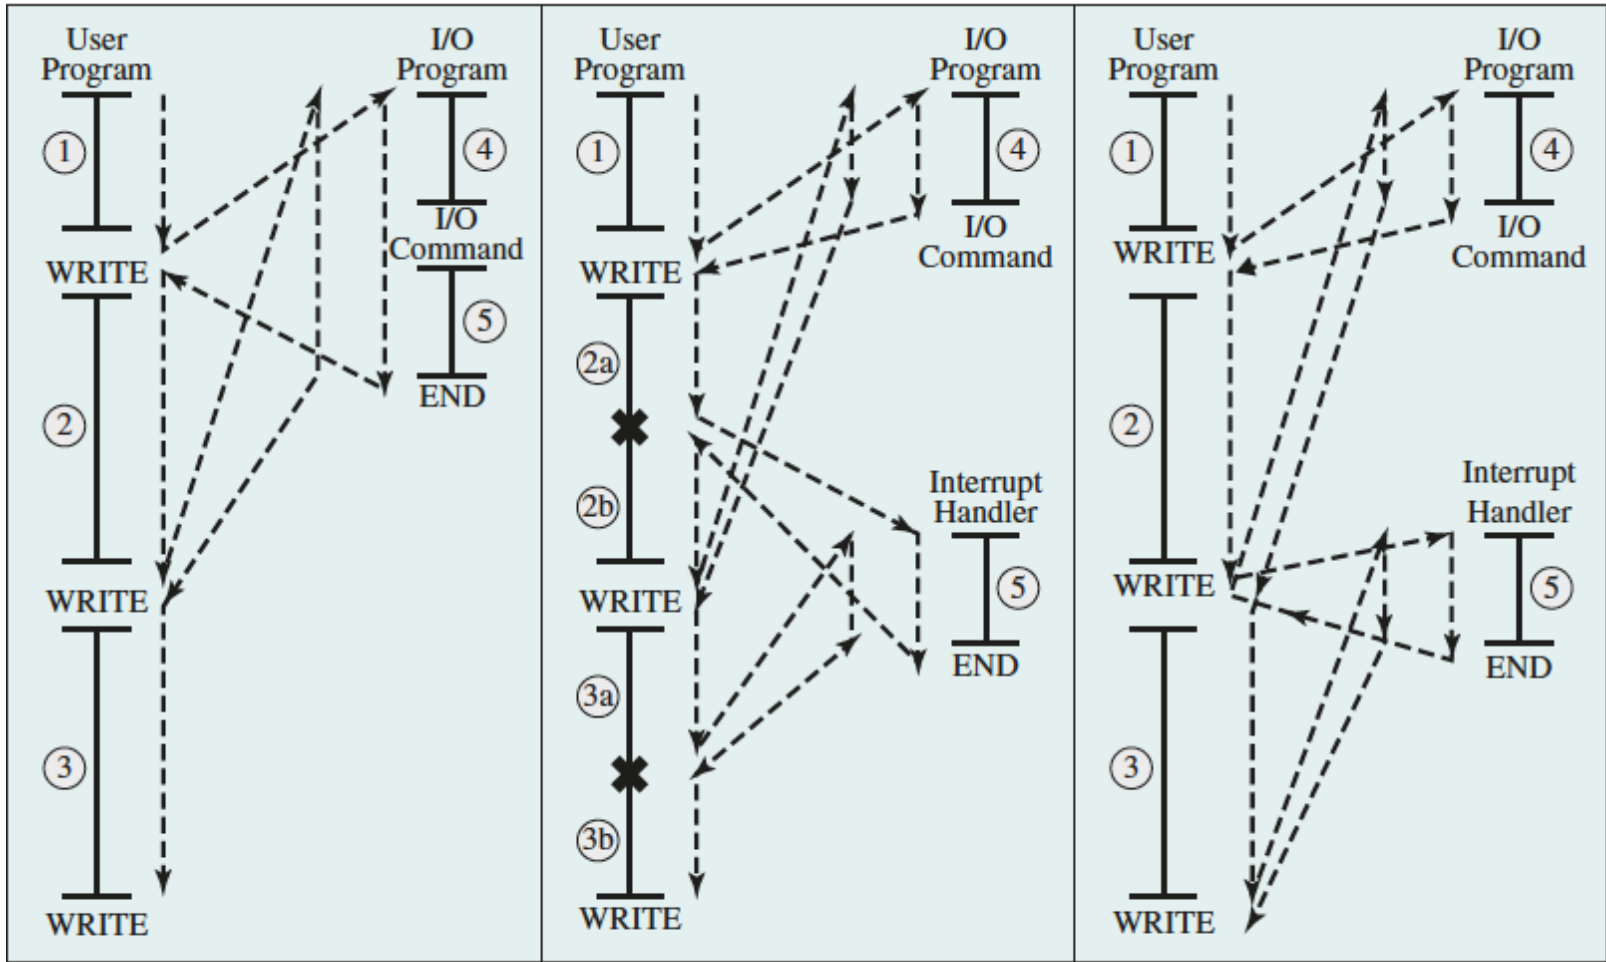
# Interrupts (cont'd)

## *Instruction Cycle without Interrupts*



## *Instruction Cycle with Interrupts*

# Interrupts (cont'd)

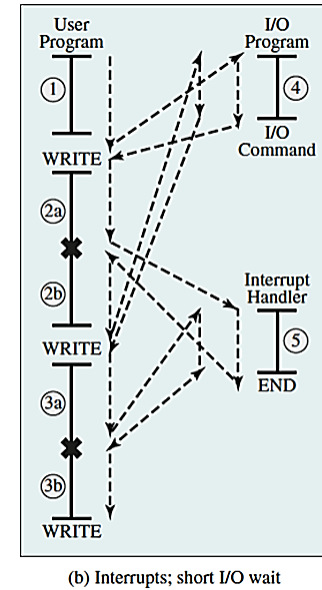
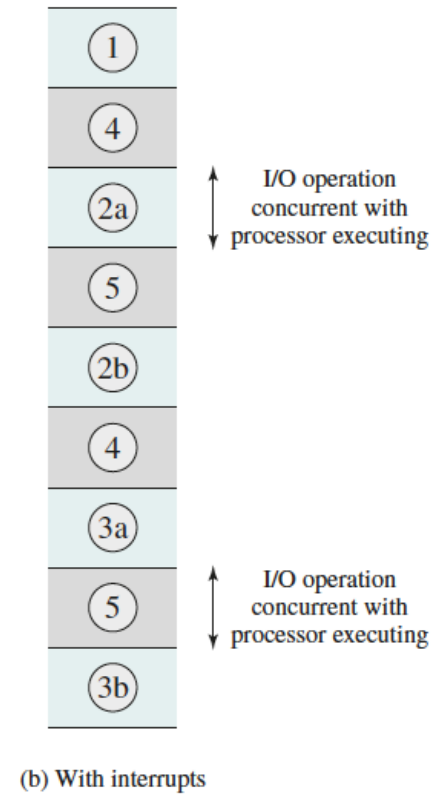



(a) No interrupts

(b) Interrupts; short I/O wait

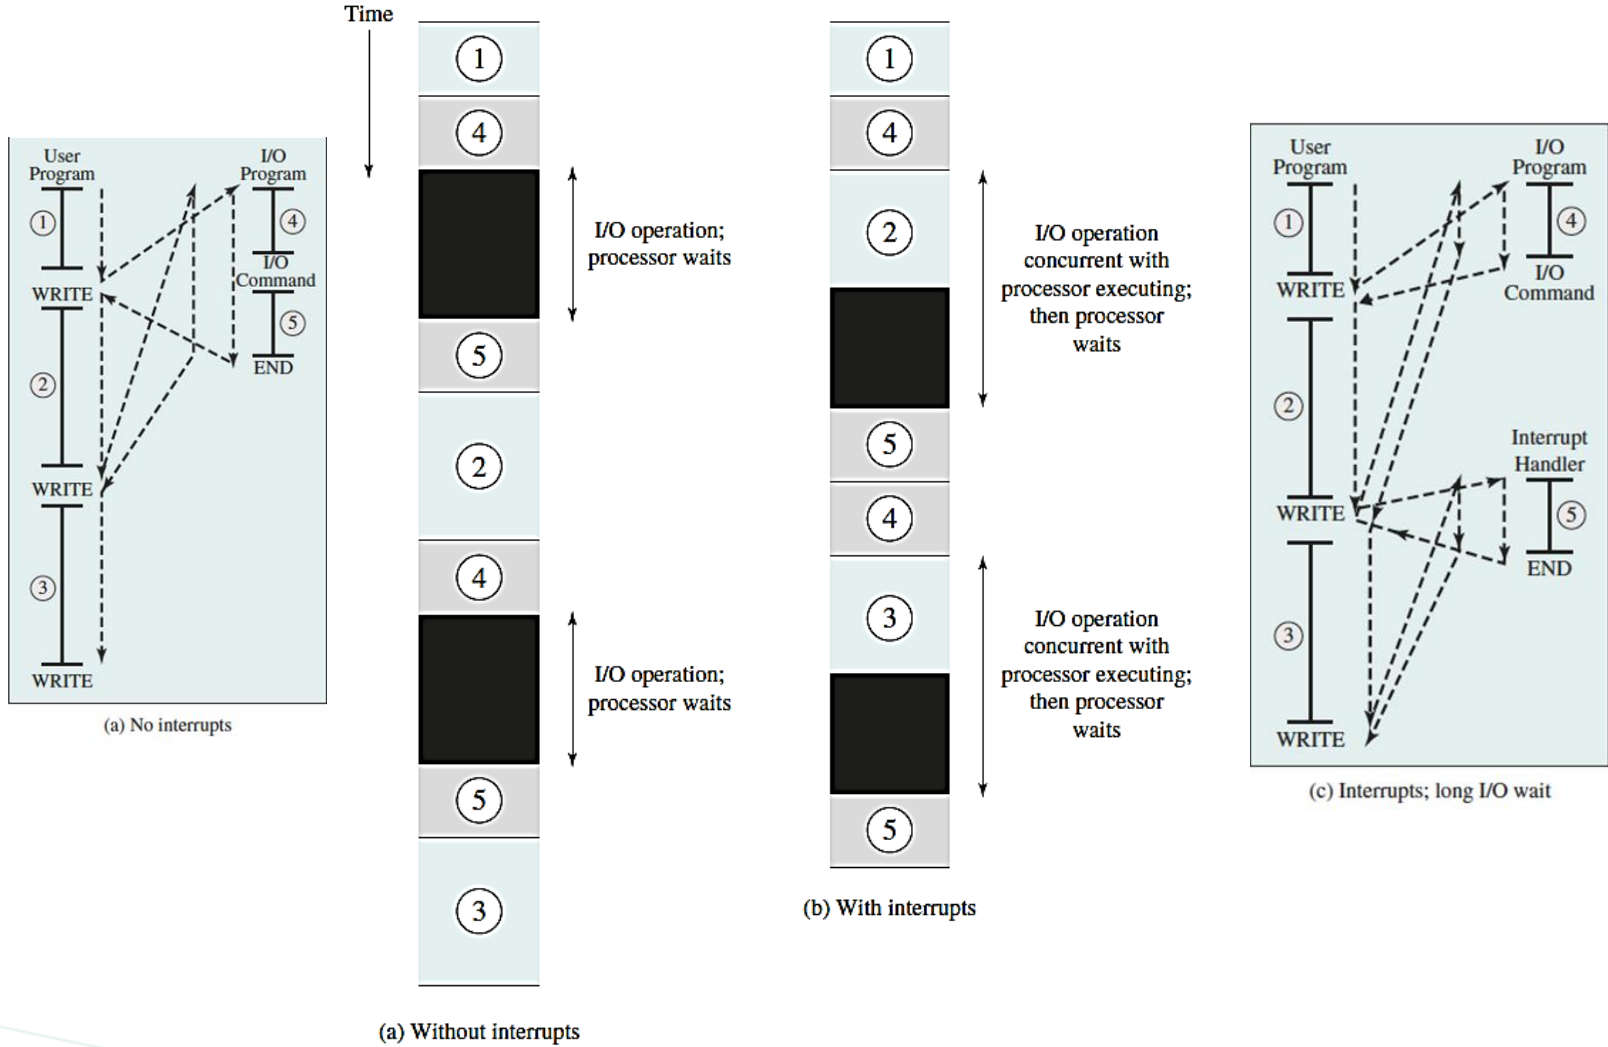
(c) Interrupts; long I/O wait

# Interrupts (cont'd)

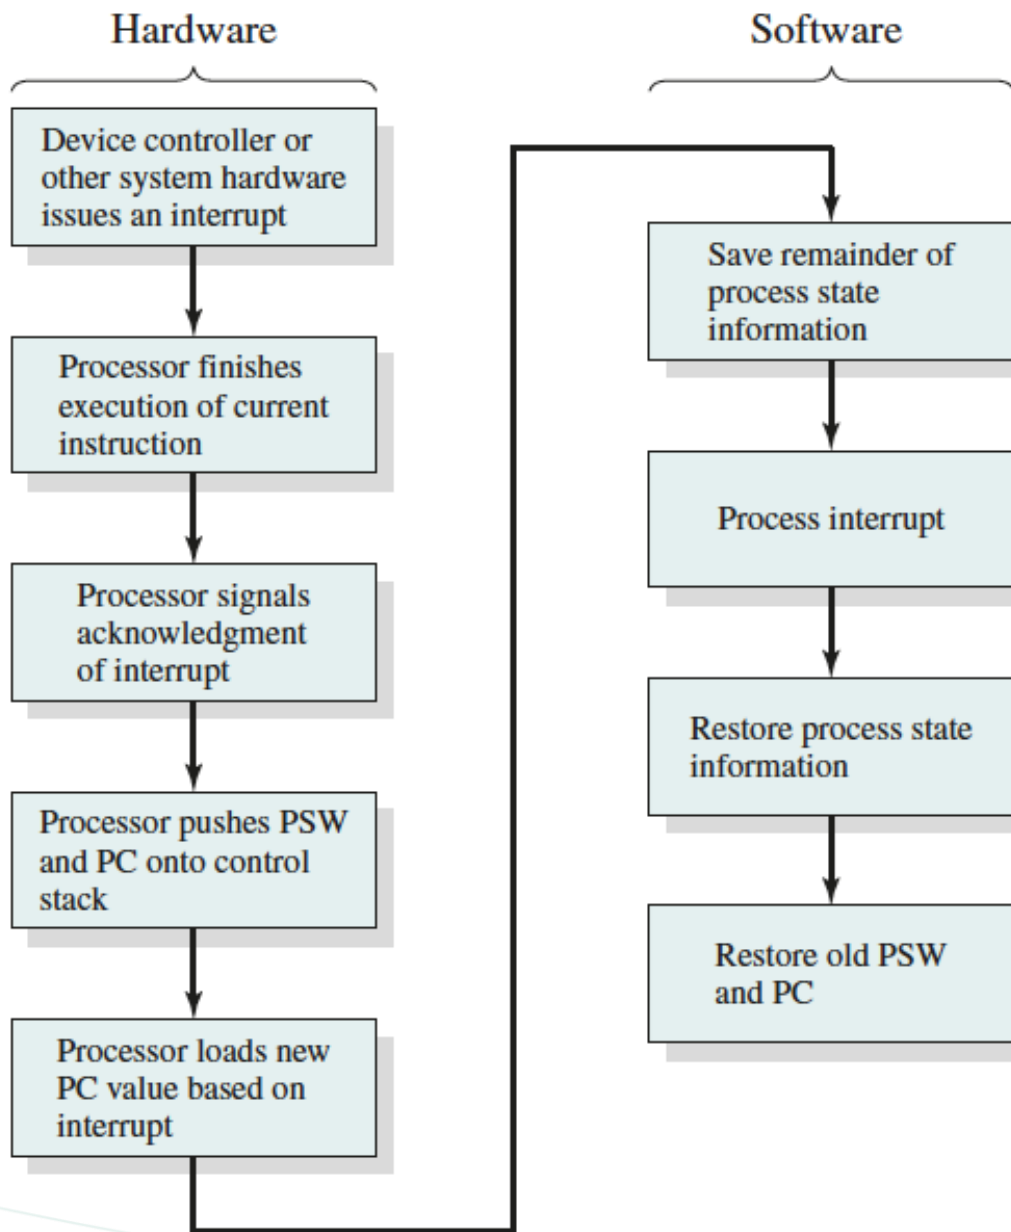




# Interrupts (cont'd)

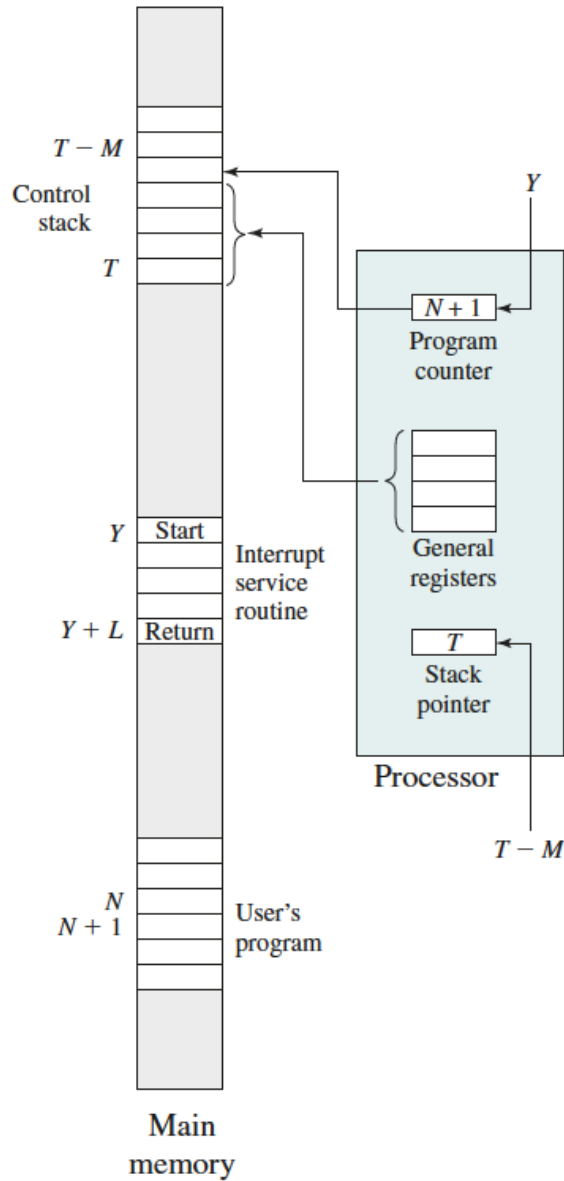


# Interrupt processing

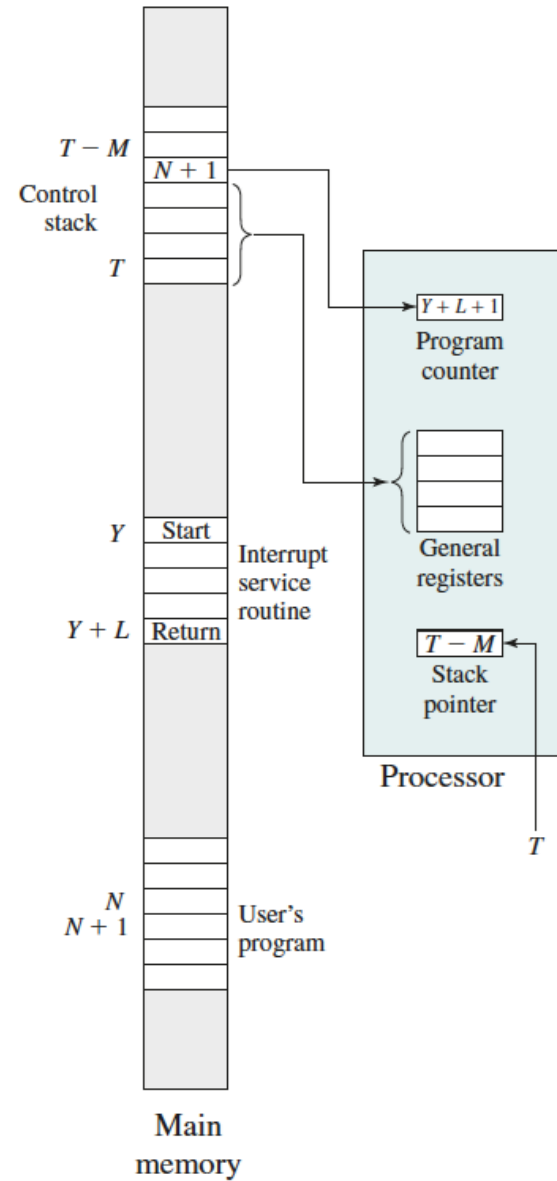


The **PSW (Program Status Word)** contains status information about the currently running process, including memory usage information, condition codes, and other status information

# Interrupt processing



(a) Interrupt occurs after instruction at location  $N$



(b) Return from interrupt


# Multiple Interrupts

## ▪ Sequential Interrupt Handling

- Disable interrupts while treating an interrupt
- Interrupts occurring while handling the current interrupt will remain pending
- Re-enable the interrupt after finishing the interrupt handling
- Disadvantages
  - Doesn't take into consideration relative priority or time-critical needs
  - Examples: communication data interrupts while printing

# Multiple Interrupts

## ▪ Sequential Interrupt Handling




# Multiple Interrupts (cont'd)

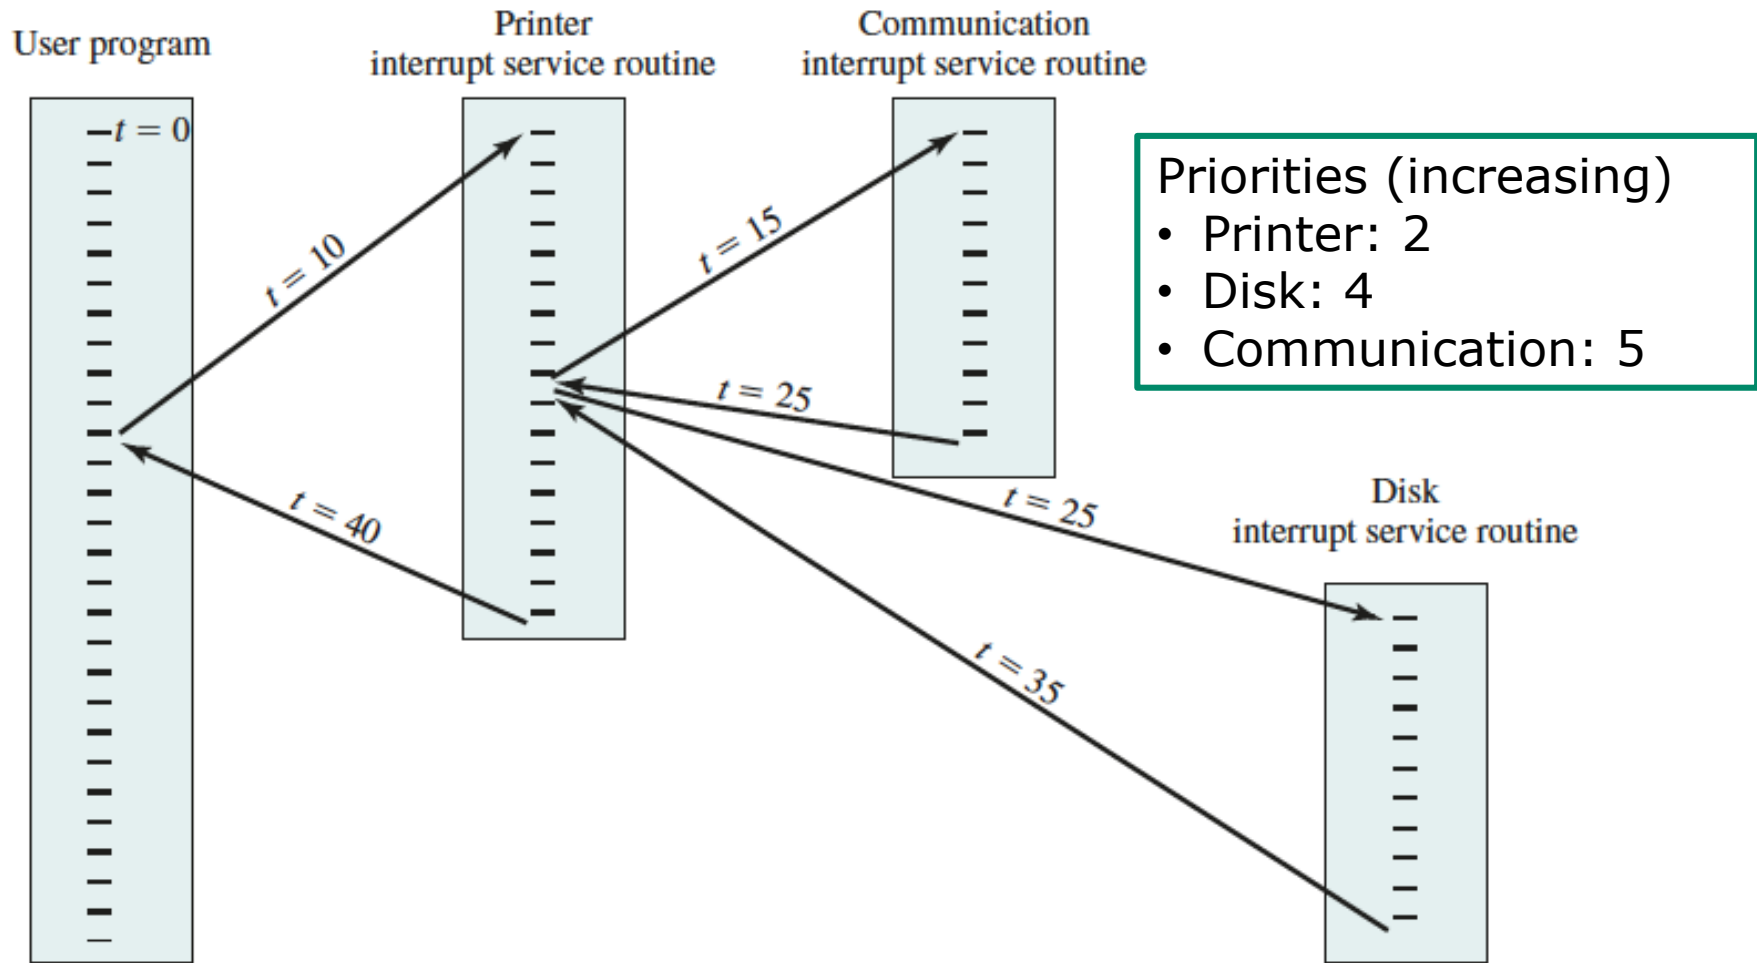
## ■ Nested interrupt processing

- Define priorities

- When an interrupt occurs while handling another interrupt, the **priority of the new interrupt** determines if it will be immediately processed or not



# Multiple Interrupts (cont'd)




# Interrupt processing

- **How interrupt handler is determined?**
  - Depending on computer architecture and OS design, there may be:
    - **Single program**
    - **One for each type of interrupt**
    - **One for each device and each type of interrupt.**
- If there is **more than one interrupt-handling routine**, the processor must determine which one to invoke.
  - Information **included in the original interrupt signal**, or
  - The **processor may have to issue a request to the device** that issued the interrupt to request the needed information.



# Interrupt processing (cont'd)

- An "**Interrupt Vector Table (IVT)**" is a data structure that associates a list of interrupt handlers with a list of interrupt requests in a table of **interrupt vectors**.
  - An entry in the interrupt vector is the address of the interrupt handler.



# Vectored vs. Non-vectored interrupt

- A **Vectored Interrupt** is where the CPU actually knows the address of the **Interrupt Service Routine (ISR)** in advance.
- All it needs is that the interrupting device sends its unique **vector** via a data bus and through its I/O interface to the CPU.
- The CPU: (1) takes this vector, (2) checks an interrupt table in memory, and then (3) carries out the correct ISR for that device.
- The vectored interrupt allows the CPU to be able to know what ISR to carry out in **software** (memory).


# Vectored vs. Non-vector interrupt

- A **Non-vector interrupt** (also called **polled interrupt**) is where the interrupting device **never sends an interrupt vector**.
- **Non-vector interrupt** has a **common ISR**, which is *common to all non-vector interrupts in the system*.
  - Address of this common ISR is known to the CPU.
- It is a specific type of *I/O interrupt* that notifies the part of the computer containing the I/O interface that a device is ready to be read or otherwise handled but does not indicate which device.
- The **interrupt controller must poll** (send a signal out to) each device to determine which one made the request.

# Memory Hierarchy

- Memory design constraints

- How much?
- How fast?
- How expensive?



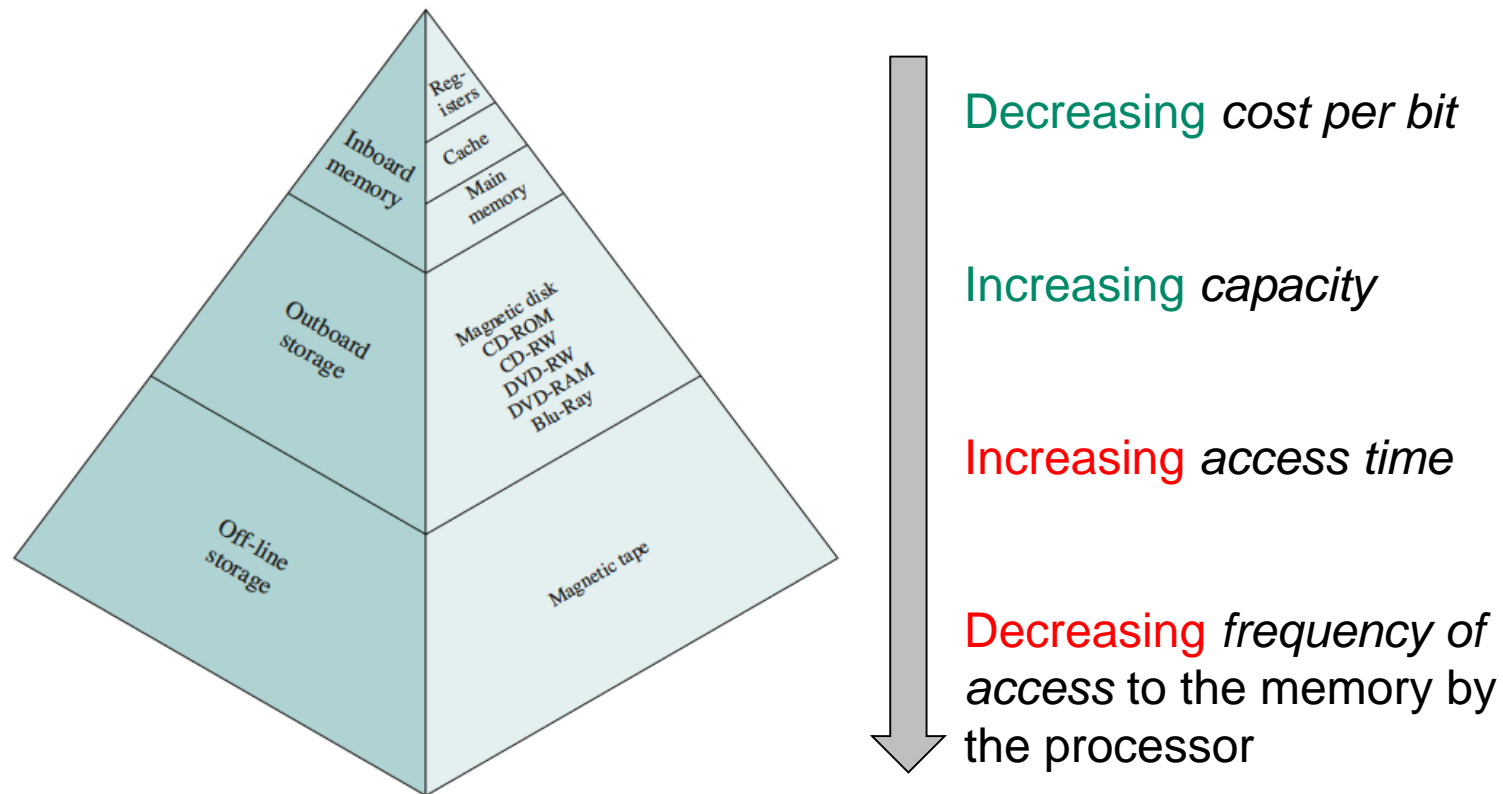
- Tradeoffs between capacity, access time and cost

- Faster access time, greater cost per bit
- Greater capacity, smaller cost per bit
- Greater capacity, slower access speed

- **How to solve the dilemma?**

# Memory Hierarchy (cont'd)

- **Solution**: Not rely on a single memory component or technology, but to employ a **memory hierarchy**



# Memory Hierarchy (cont'd)

- **Decreasing frequency of access to the memory by the processor**


- Principle of **Locality of Reference**

- During the course of execution of a program, **memory references** (instructions and data) **by the processor, tend to cluster**
- Over a long period of time, the clusters in use change, but over a short period of time, the processor is primarily working with fixed clusters of memory references

# Memory Hierarchy (cont'd)

## ■ Secondary Memory

- Auxiliary memory
- External
- Non-volatile
- Used to store program and data files



# Cache Memory: Motivation


- On all instruction cycles, **the processor accesses memory at least once, to fetch the instruction**, and often one or more additional times, to fetch operands and/or store results.
- The **rate** at which the processor can execute instructions is clearly **limited by the memory cycle time** (The time it takes to read one word from or write one word to memory).
- **This limitation has been a significant problem** because of the persistent **mismatch between processor and main memory speeds**
- **Solution**: Exploit the **principle of locality** by providing a **small, fast memory between the processor and main memory**, namely the **cache**



# Cache Memory: Principle

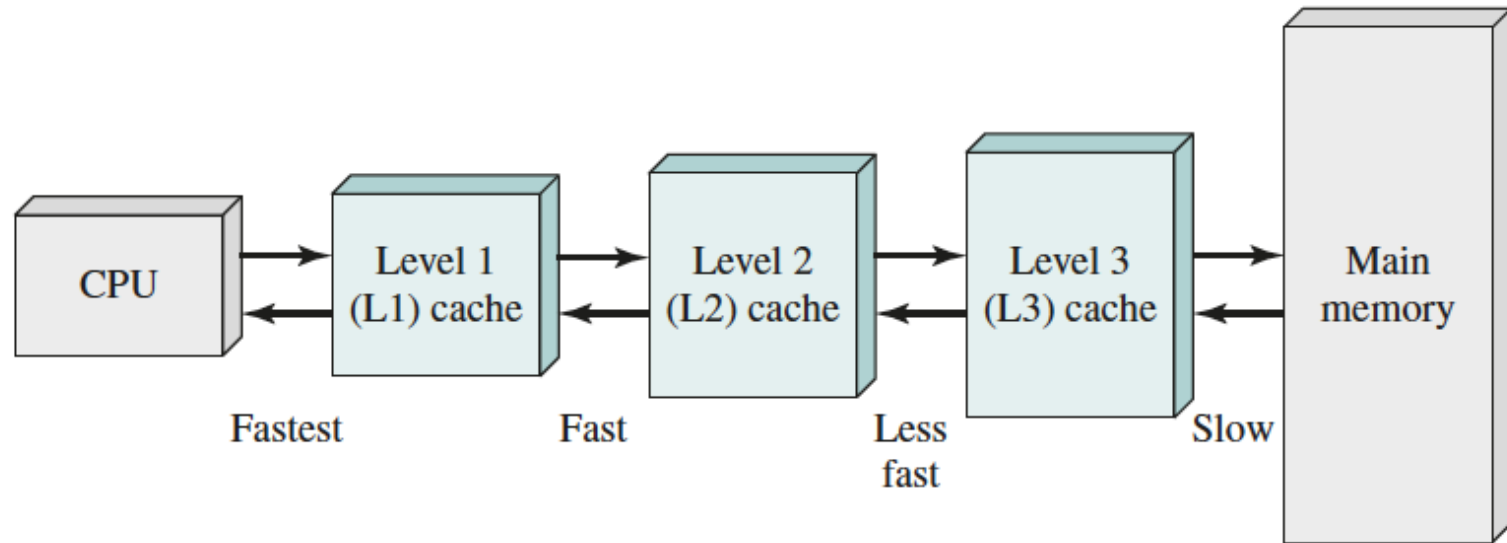
## ■ Objectives

- Provide **memory access time** approaching that of the fastest memories available
- Support a **large memory size** that has the price of less expensive types of semiconductor memories.




# Cache Memory: Principle


- Multi-level caching



# Cache Memory

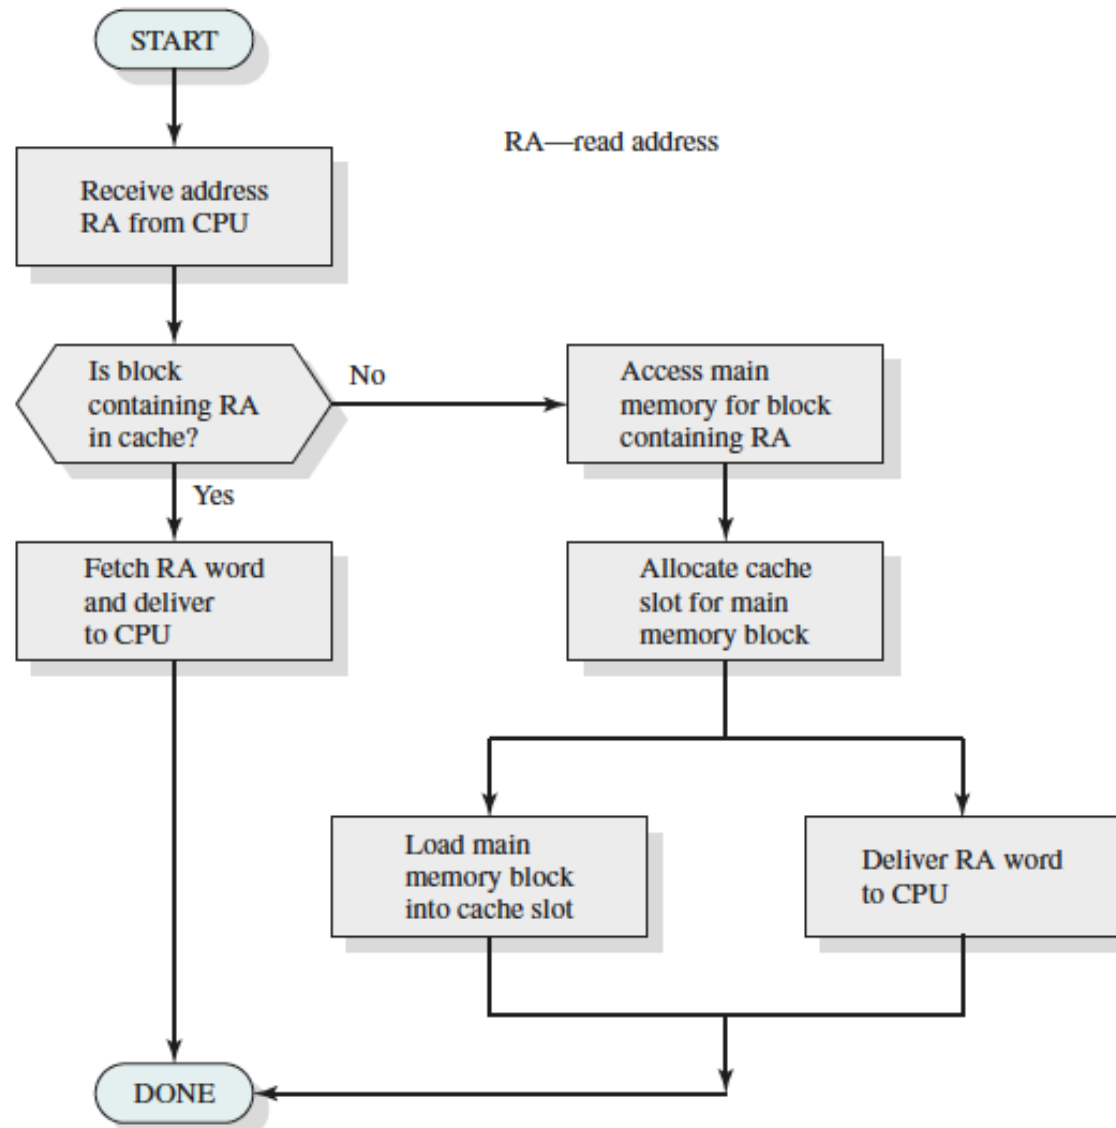


(a) Cache



(b) Main memory

# Cache Memory



# Cache Principles

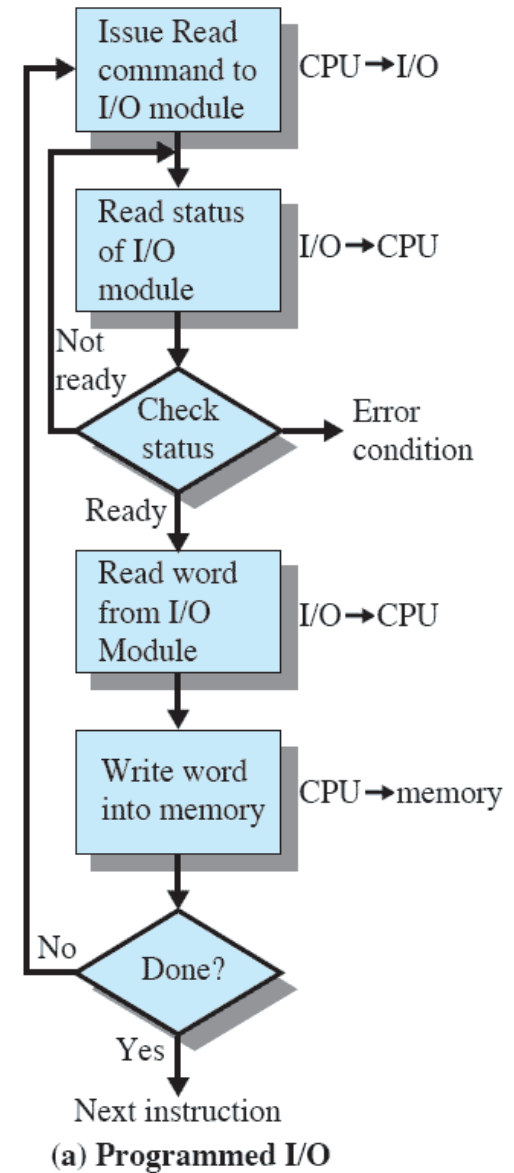
- **Cache size:** Even small caches have significant impact on performance
- **Block size:** The unit of data exchanged between cache and main memory
- **Mapping function:** Determines which cache location the block will occupy
- **Replacement algorithm:** Chooses which block to replace; e.g., **Least-recently-used (LRU)** algorithm
- **Write policy**
  - If the contents of a block in the cache are altered, then it is necessary to write it back to main memory before replacing it.
  - Dictates when the memory write operation takes place
  - Can occur every time the block is updated
  - Can occur when the block is replaced
    - Minimize write operations
    - Leave main memory in an obsolete state

# I/O Operations

- When the processor is executing a program and encounters an instruction relating to I/O, it executes that instruction by issuing a command to the appropriate I/O module
- Three possible techniques for I/O operations:
  - Programmed I/O
  - Interrupt-driven I/O
  - Direct Memory Access (DMA).

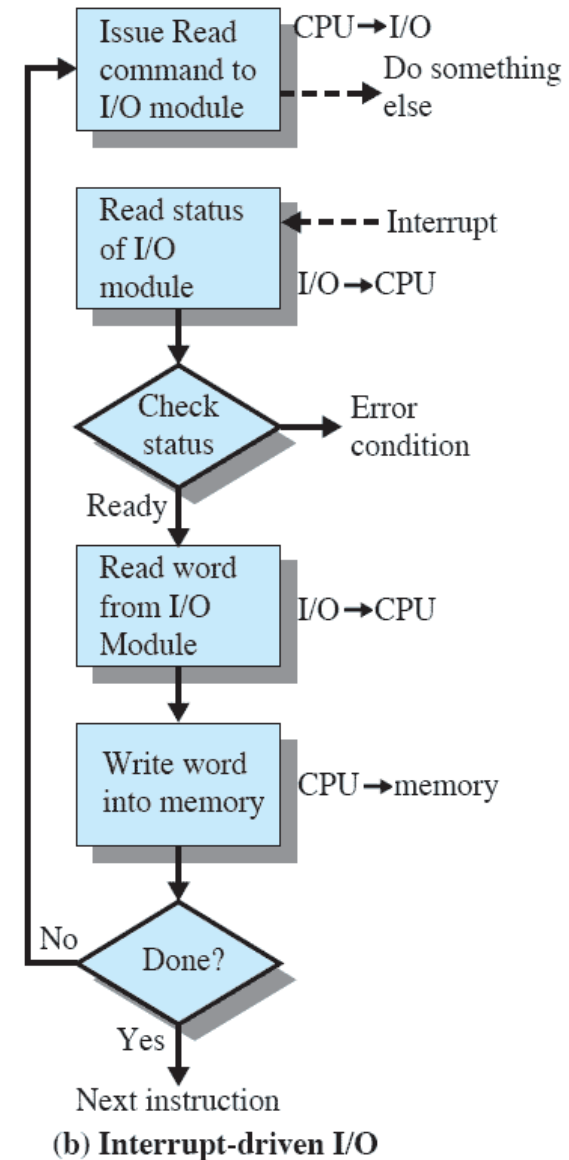
# Programmed I/O

- I/O module performs the action, not the processor
- Sets the appropriate bits in the I/O status register
- No interrupts occur
- Processor checks status until operation is complete
- Degradation of the processor performance because of long wait time



# Interrupt-Driven I/O

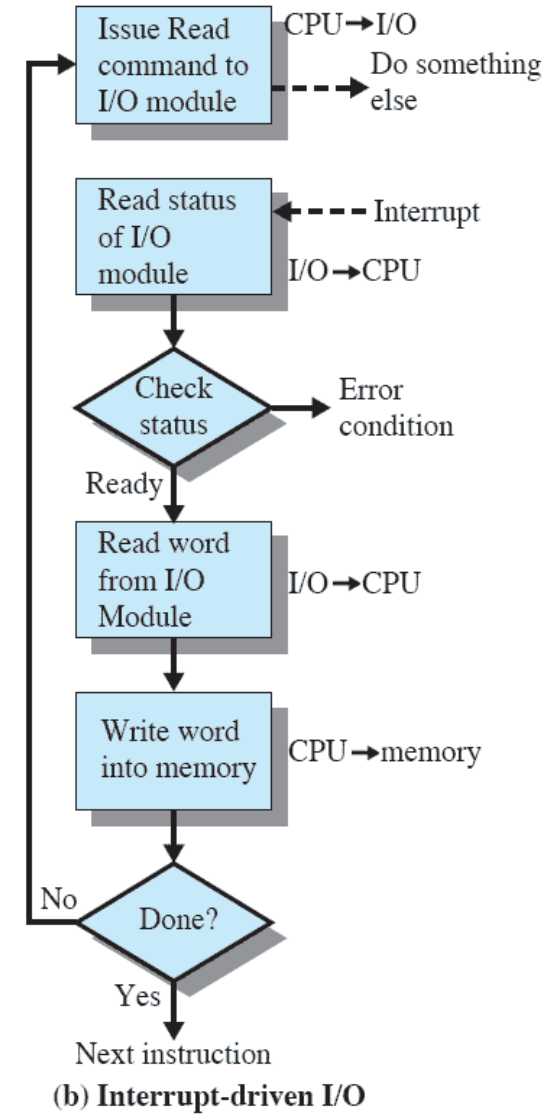
- The processor issues an I/O command to a module and then go on to do some other useful work
- **Processor is interrupted** when I/O module ready to exchange data
- Processor saves context of program executing and begins executing interrupt-handler



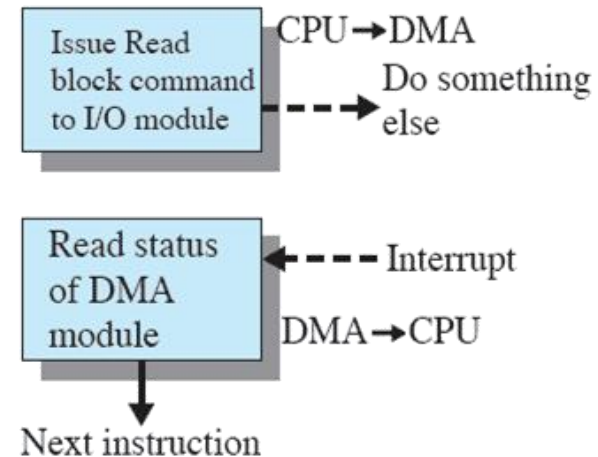
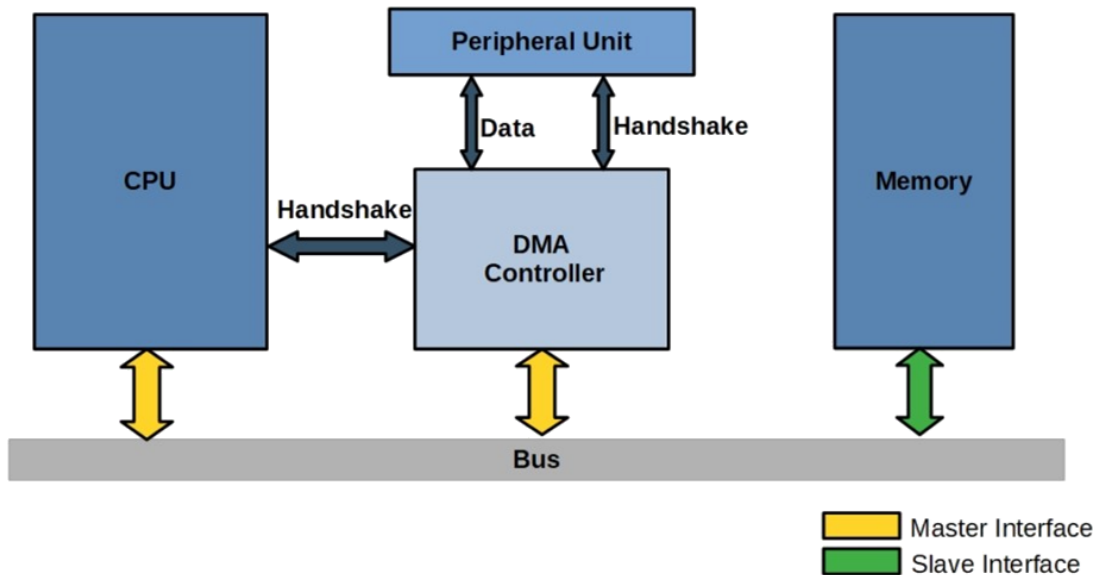


# Interrupt-Driven I/O

- No needless waiting
- Consumes a lot of processor time because every word read or written passes through the processor



# Direct Memory Access (DMA)

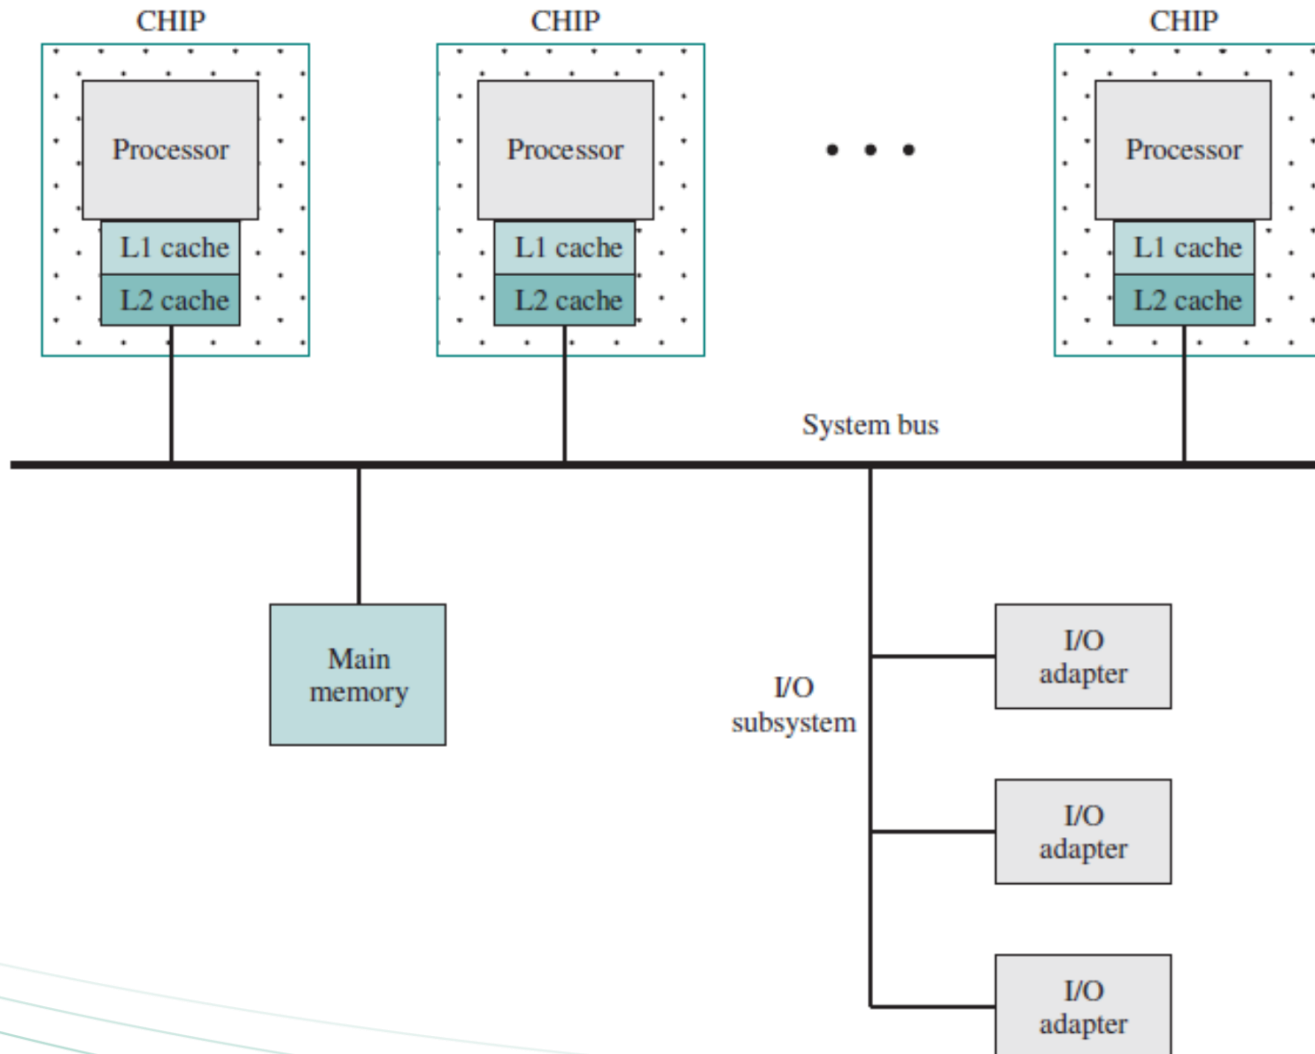


- The DMA function can be performed by a separate module on the system bus or it can be incorporated into an I/O module.
- Transfers a block of data directly to or from memory
- An interrupt is sent when the transfer is complete
- More efficient

# Multiprocessor and multicore organization

- Most popular approaches to providing parallelism by replicating processors:
  - Symmetric Multi-Processors (SMPs)
  - Multicore computers
  - Clusters

# Symmetric Multiprocessors (SMP)



# Symmetric Multiprocessors (SMP)

- A stand-alone computer system with the following characteristics:
  - Two or more similar processors of comparable capability
  - Processors share the same main memory and are interconnected by a bus or other internal connection scheme
  - Processors share access to I/O devices
  - All processors can perform the same functions
  - The system is controlled by an integrated operating system that provides interaction between processors and their programs at the job, task, file, and data element levels

# SMP Advantages

## Performance

- a system with multiple processors will yield greater performance if work can be done in parallel

## Availability

- the failure of a single processor does not halt the machine

## Incremental Growth


- An additional processor can be added to enhance performance

## Scaling

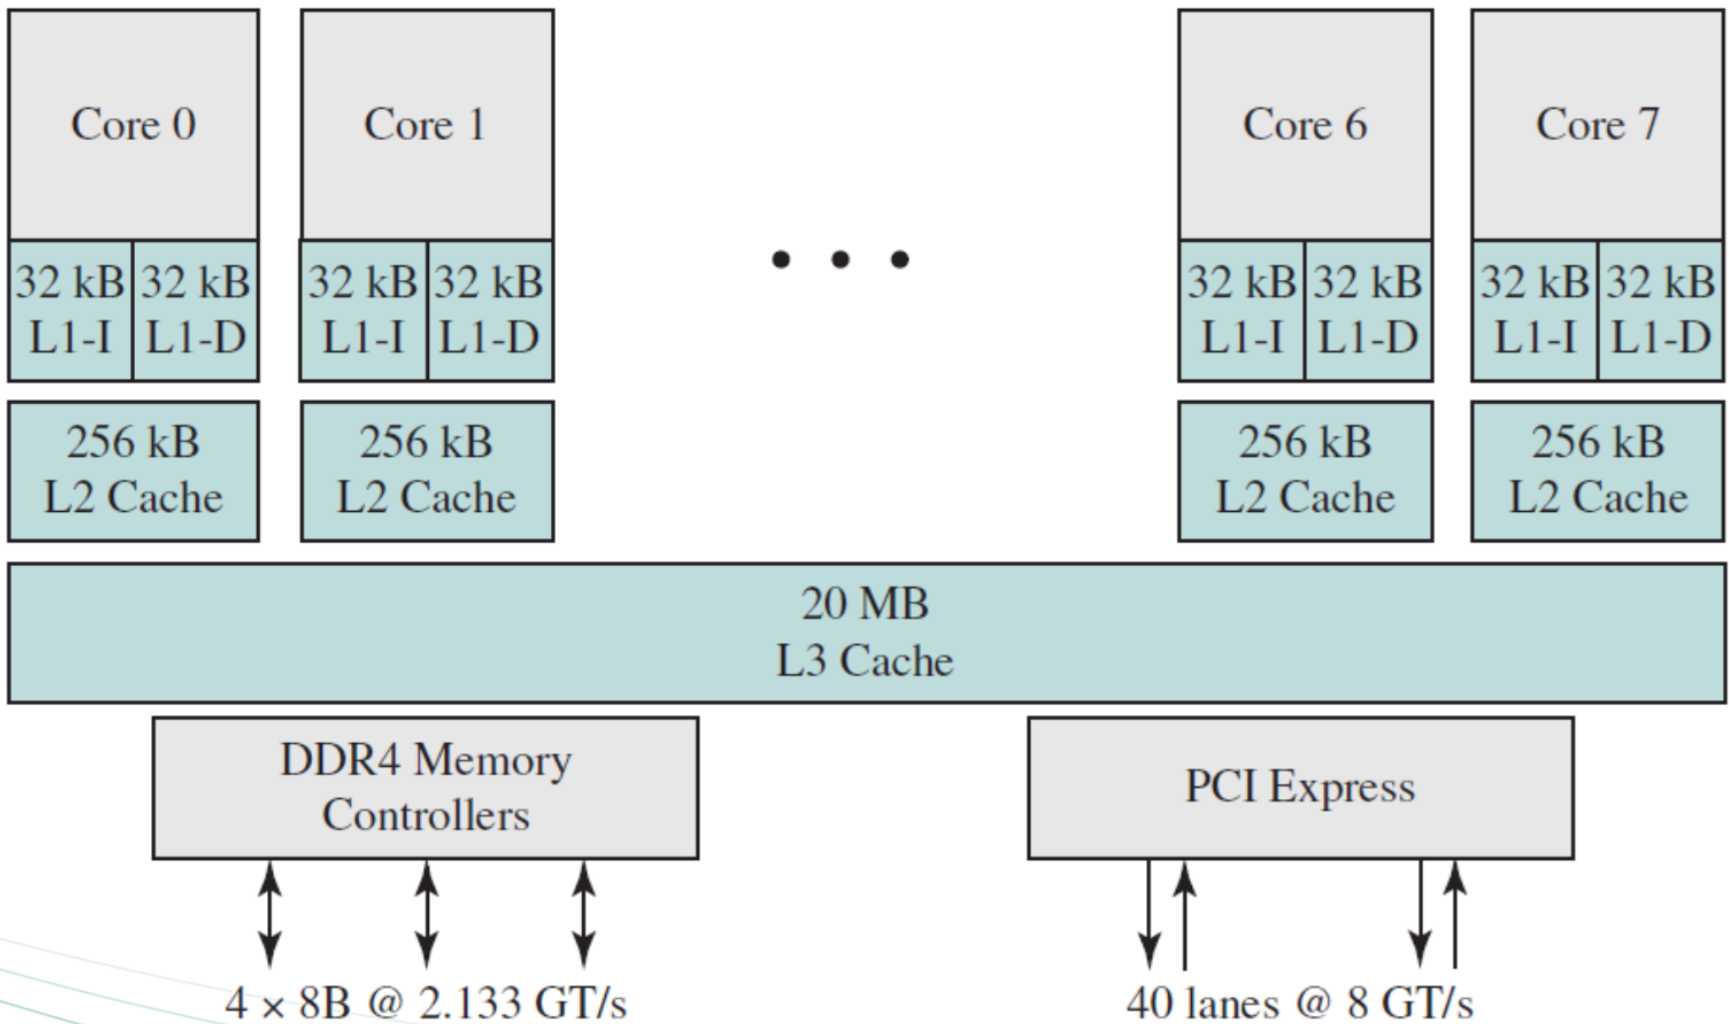
- Vendors can offer a range of products with different price and performance characteristics

# Multicore Computer

- Also known as a **chip multiprocessor**
- Combines two or more processors (**cores**) on a single piece of silicon (die)
  - Each core consists of all of the components of an independent processor
- Multicore chips also include L2 cache and in some cases L3 cache



# Multicore Computer





# Clusters

